

Joshua Chang

🌐 joshuachang.me | ✉ j79chang@uwaterloo.ca | 🔗 joshjchang | 📺 JoshChang8

EDUCATION

University of Waterloo

Sept 2021 - Apr 2026

Candidate of Bachelor of Applied Science in Systems Design Engineering, Minor in Computer Science

- Relevant Coursework: Data Structures and Algorithms, Distributed Systems, Machine Learning

EXPERIENCE

Cloud Infrastructure Engineer | The Washington Post

Washington, DC

AWS (Lambda, Kinesis, CloudWatch), PostgreSQL, Kubernetes, Docker, Terraform, Airflow, Datadog

Jan 2024 - Apr 2024

- Created an **AWS Lambda function** to **scale Kinesis data streams** to eliminate under and over utilization
- Migrated **AWS Aurora dashboards** to **Datadog**, collecting **50+** metrics of key health and performance indicators
- Created monitoring and alerting for **250+ Airflow DAGs** to catch **30+** anomalies by creating a Datadog metrics monitor provisioned using **Terraform**

Site Reliability Engineer | SAP

San Jose, CA

Docker, Kubernetes, CI/CD, Jenkins, BASH, SQL, Python.

Sep 2023 - Dec 2023

- Implemented a **Dockerized ETL** pipeline using **Kubernetes CronJobs**, ensuring reliable scheduling for data workflows
- Upgraded **Jenkins** production server and associated plugins to address security and vulnerability issues, **increasing job performance build times** and **CI/CD pipelines** by **15%**
- Created a BASH shell script to **automate** the process of **backing up SQL databases** to populate Nexus repositories, reducing the amount of manual labor spent by **85%**

Platform Engineer | Quantolio Financial Technologies

Montreal, QC

Python, PostgreSQL, Docker, Kubernetes, Heroku, Pandas

Jan 2023 - May 2023

- Containerized applications using **Docker** and orchestrated deployment on **Heroku**, ensuring accessibility for clients
- Displayed cumulative returns through dynamic graph visualizations using **Matplotlib** to improve portfolio analysis
- Performed data analysis and **feature engineering** on Bloomberg and FactSet datasets to create financial metrics

Machine Learning Researcher | University of Waterloo

Waterloo, ON

Python, MySQL, PyTorch, Pandas, AWS (S3), Pose Estimation, CNN Models

Sep 2022 - Dec 2022

- Utilize **YOLOv5** for object detection to accurately identify baseball pitchers, ensuring precise tracking and recognition
- Integrated a **3D Pose Estimation model** for biomechanical analysis to improve performance and aid in injury prevention
- Aggregate and clean Hawkeye baseball data from **SQL databases** and **AWS S3** storages to train **CNN models**

PROJECTS

Spotify Song Recommender 🎵

Python, Pandas, NumPy, Jupyter

- Deployed a webapp with **4,500+ users**, analyzing Spotify playlists to deliver personalized song recommendations
- Collected and processed dataset of **1M+ songs** using **Pandas** and **NumPy** to construct **item-feature matrices**
- Implemented **multiprocessing** to enhance efficiency of CPU operations, resulting in **35%** reduction in computation time
- Orchestrated containerization of web application using **Docker** and enabled **automatic deployment** on Google Cloud Run with Google Build triggers for **continuous integration**

SKILLS

Languages: Python, SQL, Java, C++, Git, BASH

Tools and Technologies: AWS, Kubernetes, Docker, Helm, Terraform, Airflow, SQL, Datadog, CUDA

Technical Skills: Cloud Infrastructure, MLOps, CI/CD, Distributed Systems, ETL Pipelines, API Development